

Relational data

Four main types of operations with two tables

- **Binding**, which simply stacks tables on top of or beside each other
- **Mutating joins**, which add new variables to one data frame from matching observations in another.
- **Filtering joins**, which filter observations from one data frame based on whether or not they match an observation in the other table.
- **Set operations**, which treat observations as if they were set elements.

Keys

- A variable (**or set of variables**) that uniquely identifies an observation
 - A **primary key** uniquely identifies an observation in its own table [can be a set of variables]. For example, `planes$tailnum` is a primary key because it uniquely identifies each plane in the `planes` table.
 - A **foreign key** uniquely identifies an observation in another table [can be a set of variables]. For example, the `flights$tailnum` is a foreign key because it appears in the `flights` table where it matches each flight to a unique plane.

Relations

- Typically one-to-many
 - Each flight has one plane, but each plane has many flights
- Can also be many-to-many
 - Each airline flies to many airports; each airport hosts many airlines
- Can also be one-to-one

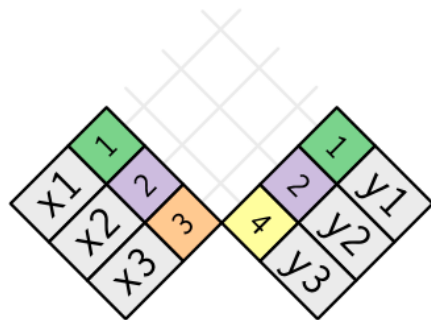
Understanding joins

x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y3

Understanding joins

x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y3

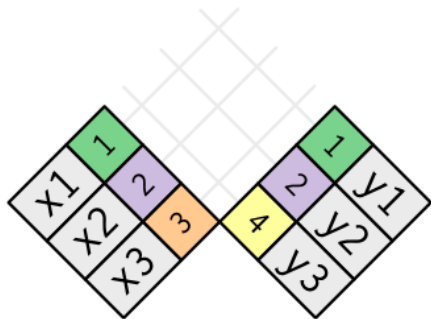
Each potential match



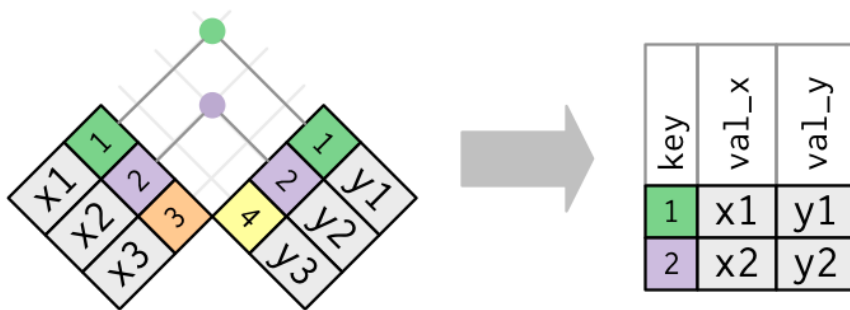
Understanding joins

x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y3

Each potential match




Number of actual matches




Inner join: Unmatched rows are not included in the output

Outer joins


Keeps all observations in x



Keeps all observations in y



Keeps all observations in x and y



Combine Data Sets

a		b	
x1	x2	x1	x3
A	1	A	T
B	2	B	F
C	3	D	T

+

=

Mutating Joins

x1	x2	x3
A	1	T
B	2	F
C	3	NA

dplyr::left_join(a, b, by = "x1")

Join matching rows from b to a.

x1	x3	x2
A	T	1
B	F	2
D	T	NA

dplyr::right_join(a, b, by = "x1")

Join matching rows from a to b.

x1	x2	x3
A	1	T
B	2	F

dplyr::inner_join(a, b, by = "x1")

Join data. Retain only rows in both sets.

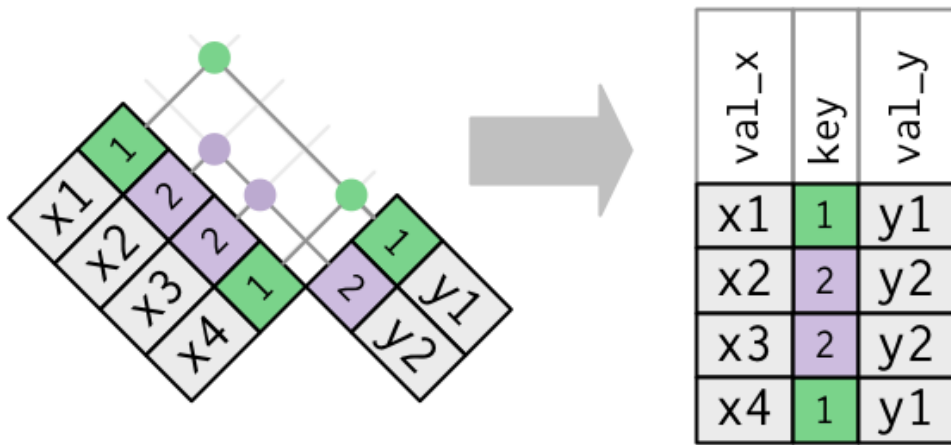
x1	x2	x3
A	1	T
B	2	F
C	3	NA
D	NA	T

dplyr::full_join(a, b, by = "x1")

Join data. Retain all values, all rows.

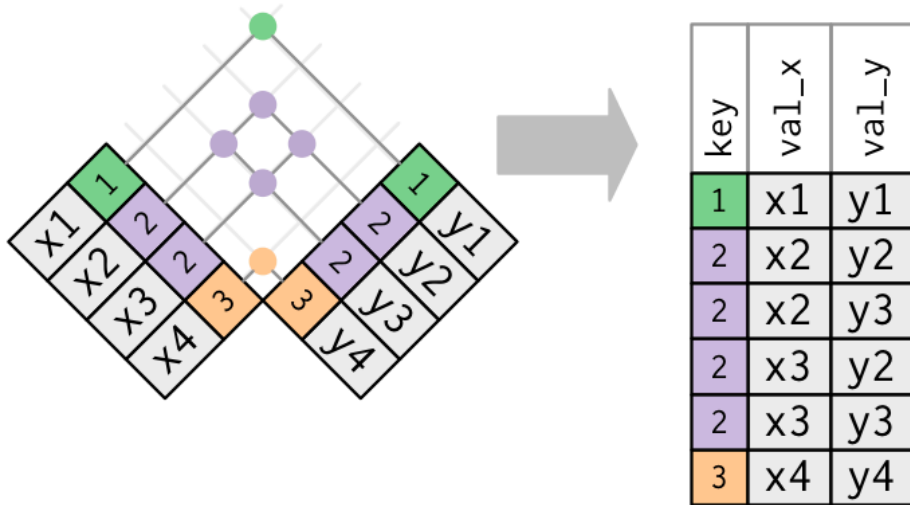
Duplicate keys

One table has duplicate keys (typically a one-to-many relationship)
e.g. "dest" in the flights tibble



Duplicate keys

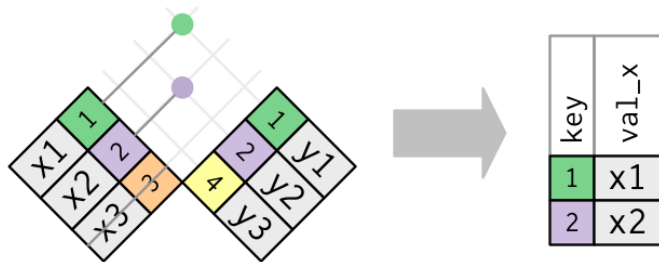
Both tables have duplicate keys (typically an error)




Filtering

- `semi_join(x, y)` **keeps** all observations in x that have a match in y.
- `anti_join(x, y)` **drops** all observations in x that have a match in y.

Semi-join



Anti-join



Exercises

- Filter flights to only show flights with planes that have flown at least 100 flights
- Combine `fueleconomy::vehicles` and `fueleconomy::common` to find only the records for the most common models

Join problems – how to troubleshoot

- Start by identifying the variables that form the primary key in each table based on your understanding of the data
- Check that none of the variables in the primary key are missing. If a value is missing then it can't identify an observation!
- Check that your foreign keys match primary keys in another table. The best way to do this is with an `anti_join()`

Factors

Exercises

- Explore the distribution of rincome (reported income). What makes the default bar chart hard to understand? How could you improve the plot?
- What is the most common relig in this survey? What's the most common partyid?